

# Assignment 3: Sentiment Analysis with Convolutional Neural Networks

Filip Stefaniuk  
filipste@student.matnat.uio.no

October 25, 2018

## Introduction

In this assignment, I have implemented Convolutional Neural Network classifier based on Kim<sup>1</sup> 2014 and Ye Zang<sup>2</sup> 2015 work. I haven't achieved the accuracies listed in papers, but I suspect this is due to the fact that the dataset we have in this assignment is different. The authors of the mentioned papers used also the **phrases** part of the SST-2 dataset for training.

I have implemented two scripts:

- `train_model.py` - used for training
- `eval_on_test.py` - used for evaluation

Additionally I have several functions splitted across the python modules:

- `data.py` - data preprocessing
- `emb.py` - loading embeddings and creation of embeddings layers
- `model.py` - functions for creating the models

All the experiments were run on abel with the scripts provided in the `scripts` directory. Source code, scripts, logs from experiments, results as json files, notebooks used to present the data and sourcecode of this report available in github repository<sup>3</sup>

## 1 Baseline sentiment analisys

I have build CNN model with hyperparameters listed in the assignment, namely:

1. categorical cross-entropy loss
2. pre-trained 300-D *word2vec* embeddings trained on Google News corpus, I used original version downloaded from the web.
3. sequence of static word embeddings
4. window sizes of 3, 4 and 5
5. 100-D filters
6. ReLU activation function
7. global 1-max pooling
8. dropout with dropout rate of 0.5

I have not restricted the vocabulary size resulting in having 13295 words in vocabulary. I have padded and truncated sentences to length of 50. As a result I received architecture listed below.

---

<sup>1</sup>Yoon Kim 2014

<sup>2</sup>Ye Zang et al. 2015

<sup>3</sup><https://github.com/filipste/INF5820-Assignment3>

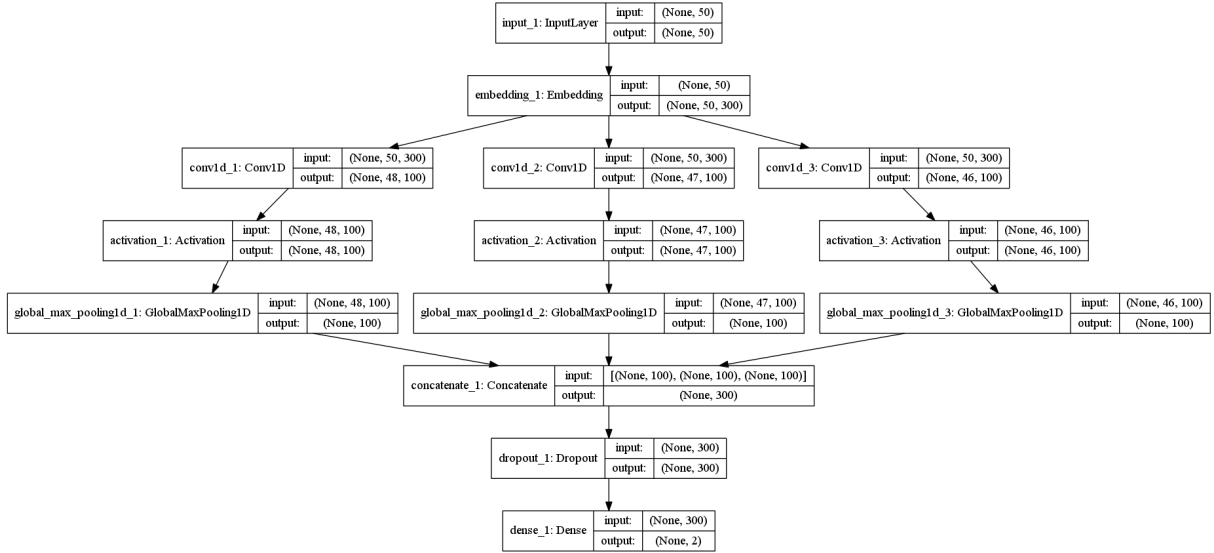


Figure 1: Baseline model architecture

I have trained the model using batch size of **256** and **adadelat** optimizer (same as in the papers).

## 2 Controlling the randomness

I have trained the model 10 times both with unset seed and seed set to **123**. I got the following results:

|         | min   | max   | mean  | std   |
|---------|-------|-------|-------|-------|
| no seed | 0.767 | 0.783 | 0.777 | 0.006 |
| seed    | 0.776 | 0.799 | 0.789 | 0.007 |

## 3 Sensitivity analysis: tuning hyperparameters

When testing hyperparameters I was changing value of only one of them keeping the rest as in a baseline model. The authors of the *A Sensitivity Analysis of (and Practitioners' Guide to) Convolutional Neural Networks for Sentence Classification* concluded that almost always 1-max pooling strategy outperformed all the other approaches and dropout is the best regularization term. Because of that I decided that those two hyperparameters are not interesting to test and experimented with all the others. I set the same seed for all the experiments.

### 3.1 Activation function

I have tested multiple activation functions in convolutional layers. The *sigmoid* and *tanh* activations performed poorly. I tried different variations of *ReLU* activation namely *eLU* and *leakyReLU* but there was no improvement. Seems like the *ReLU* is the best activation function to use in this case, what is similar to the conclusions of authors of *A Sensitivity Analysis...* paper.

| activation | accuracy        |
|------------|-----------------|
| tanh       | 0.768485        |
| sigmoid    | 0.665455        |
| elu        | 0.780606        |
| relu       | <b>0.793939</b> |
| leakyrelu  | 0.789091        |

Figure 2: Accuracy with different activations

### 3.2 Convolutional window size

I have tested the same window sizes as in the previously mentioned paper, it seems that in this case there is not much of a difference when using different window sizes. Turns out that the best setting was either using the 3-4-5, 1-3-5-7 or simply 3 window size. Seems like window size of 3 is enough to capture necessary information.

### 3.3 Number of filters per window size

Experimenting with the filter size, yielded results that the optimal size of the filter is 200.

| window sizes | accuracy        |
|--------------|-----------------|
| 1            | 0.787879        |
| 1-3-5-7      | 0.792727        |
| 14-15-16     | 0.753939        |
| 2-3-4        | 0.786667        |
| 2-3-4-5      | 0.773333        |
| 3            | <b>0.796364</b> |
| 3-4-5        | 0.790303        |
| 4-5-6        | 0.780606        |
| 5            | 0.773333        |
| 6-7-8-9      | 0.761212        |
| 7-8-9        | 0.761212        |

(a) Accuracy with different window sizes

|   | filter_size | accuracy        |
|---|-------------|-----------------|
| 0 | 50          | 0.772121        |
| 1 | 100         | 0.786667        |
| 2 | 200         | <b>0.791515</b> |
| 3 | 300         | 0.790303        |
| 4 | 500         | 0.772121        |
| 5 | 1000        | 0.785455        |

(b) Accuracy with different filter sizes

### 3.4 Raw, lemmatized and POS tagged data

I have tried using POS tagged data, but it seems that the word embedding model with POS tags is not good. Using Google News model (*1.zip*) there was **9001** out of vocabulary tokens, so the model had very poor results.

To compare results when using lemmatized and non-lemmatized data I used models trained on Wikipedia and Gigaword (*17.zip*, *18.zip*, *19.zip* *20.zip*). There was no clear improvement when using lemmatized data.

|          | f1    | precision | recall |
|----------|-------|-----------|--------|
| negative | 0.679 | 0.659     | 0.701  |
| positive | 0.666 | 0.688     | 0.645  |
| average  | 0.673 | 0.673     | 0.673  |

(a) Metrics when using model 1.zip

|                       | raw   | lemmatized |
|-----------------------|-------|------------|
| wiki+Gigaword (w2v)   | 0.655 | 0.661      |
| wiki+Gigaword (glove) | 0.753 | 0.765      |

(b) Comparison of using raw and lemmatized data

### 3.5 Multiple channels

I have tested model with multiple channels, one with static and the second with no-static word embeddings as in Kim 2014. Results are presented in section about influence of word embeddings.

## 4 Testing the influence of word embeddings

I have tested multiple word embeddings in 3 modes: **static**, **non-static** and **multichannel**. Fine tuning helps achieve better results, but the best scores were achieved when using **multichannel** mode. Fasttext turned out to be the best model.

|                          | static       | non-static   | multichannel |
|--------------------------|--------------|--------------|--------------|
| wiki+Gigaword (w2v)      | 0.655        | 0.659        | 0.704        |
| wiki+Gigaword (glove)    | 0.753        | 0.764        | 0.759        |
| GoogleNews (w2v)         | 0.790        | 0.795        | 0.790        |
| CommonCrawl 840B (glove) | 0.790        | 0.794        | 0.804        |
| wiki-news (fasttext)     | <b>0.796</b> | <b>0.801</b> | <b>0.818</b> |

Figure 5: Comparison of different word embeddings.

### 4.1 Inferring vectors for OOV words

Inferring OOV words further improved the results. Model trained with these embeddings achieved accuracy **0.818** in **static** mode.

|          | f1    | precision | recall |
|----------|-------|-----------|--------|
| negative | 0.812 | 0.831     | 0.794  |
| positive | 0.824 | 0.807     | 0.842  |
| average  | 0.818 | 0.819     | 0.818  |

Figure 6: Results after inferring OOV words

### 4.2 Influence of vocabulary size

I have tested how changing the vocabulary size changes the number of the oov tokens. Experiments were performed on Google News w2v.




Figure 7: Results after inferring OOV words